

The performance of the GPU-accelerated ELPA2 solver is tested on the Ascent computer at Oak Ridge National Laboratory. Each node of Ascent has 2 IBM POWER9 CPUs with 42 cores in total, and 6 NVIDIA Volta GV100 GPUs. As Fig. 1 shows, the bottlenecks seem to be the first and fourth steps.

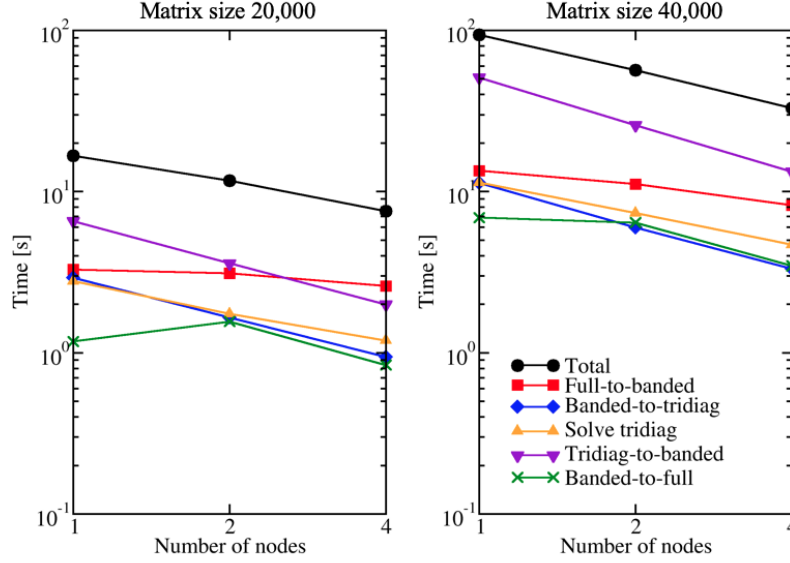


Figure 1: Time to solution of the five computational steps of ELPA2 with respect to the number of nodes. Matrix sizes are $20,000 \times 20,000$ (left) and $40,000 \times 40,000$ (right). Tests are performed on the Ascent computer at the Oak Ridge National Laboratory. Each node consists of 2 IBM POWER9 CPUs and 6 NVIDIA Volta GV100 GPUs. The nodes are fully exploited by running 42 MPI tasks and 6 GPUs per node.

This document describes a CUDA kernel for applying a series of Householder transformations to a matrix, which is the main computational task in the fourth step of ELPA2. Mathematically this is as simple as

$$Y = H_1 H_2 \cdots H_n X. \quad (1)$$

H_i is a Householder transformation matrix taking the form

$$H_i = I - \tau_i v_i v_i^*, \quad (2)$$

with τ_i being a scalar, and v_i a vector. Since a Householder matrix can be completely defined by τ_i and v_i , there is no need to explicitly construct and store H_i . Instead, only τ_i and v_i are stored.

In an actual ELPA2 calculation with MPI, each MPI process applies n Householder transformations to the local portion of the eigenvector matrix X . The length of each Householder vector v_i is b , which is always equal to the semi-bandwidth nbw of the banded matrix stage in ELPA2. The dimension of the local eigenvector matrix is $N_R \times N_C$, where the number of rows $N_R = b + n - 1$, and the number of columns N_C is a user-specified parameter (usually 1,024).

The n Householder transformations must be transformed one after another. This is done in a loop:

```
do i = n, -1, 1
  apply the nth transformation
end do
```

Fig. 2 shows the data layout of the Householder transformation CUDA kernel, using the n^{th} and $(n-1)^{th}$ iterations as an example. Here $b = 4$, $n = 4$, $N_R = b + n - 1 = 7$, and $N_C = 6$. In the n^{th} iteration,

the n^{th} Householder vector is applied to the eigenvector matrix \mathbf{X} . Note that the n^{th} \mathbf{v} only alters the last 4 rows of \mathbf{X} (yellow part in Fig. 2). This is a general property of Householder transformations. The kernel is launched with N_C blocks and b threads, such that each block works on one column of the eigenvector matrix, whereas each thread within a block in turn works on one element of the eigenvector matrix. For each block, the task is to compute

$$\mathbf{H}\mathbf{x} = (\mathbf{I} - \tau\mathbf{v}\mathbf{v}^*)\mathbf{x} = \mathbf{x} - \tau\mathbf{v}(\mathbf{v}^*\mathbf{x}). \quad (3)$$

Again, the scalar τ and vector \mathbf{v} are used to define a Householder transformation. The vector \mathbf{x} is a sub-vector of the eigenvector matrix \mathbf{X} . The lengths of \mathbf{v} and \mathbf{x} are both equal to b . Eq. 3 is computed in two steps. First, the dot product $d = \mathbf{v}^*\mathbf{x}$ is computed by multiplying the corresponding element of \mathbf{v} and \mathbf{x} on each thread, then performing a parallel reduction within a block. Second, the vector \mathbf{x} is updated by $\mathbf{x} - \tau\mathbf{v}(d)$, where the threads in a block can work independently from each other.

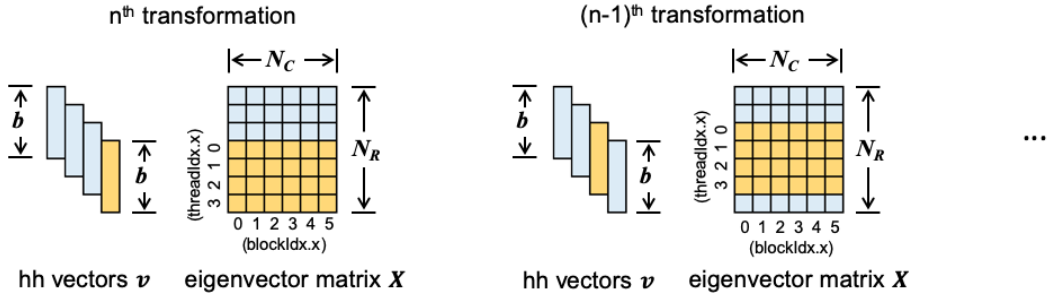


Figure 2: Workflow of the Householder transformation CUDA kernel. $b = 4$, $n = 4$, $N_R = b + n - 1 = 7$, and $N_C = 6$. A block works independently on a column of the eigenvector matrix. A thread works on one element of the eigenvector matrix. From the n^{th} iteration to the $(n - 1)^{th}$ iteration, the working part of the eigenvector matrix is shifted upward by one element.

When the n^{th} transformation is finished, the last row of the eigenvector matrix \mathbf{X} will no longer be updated. Therefore, within every block, the thread that currently holds the last row of \mathbf{X} writes the value to the result matrix. Then in the $(n - 1)^{th}$ transformation, thread t , $t > 0$, takes the value of thread $t - 1$, whereas thread $t = 0$ reads in a new row from the eigenvector matrix \mathbf{X} . The computation of Eq. 3 can proceed as before. The kernel finishes when all the n Householder transformations are applied to the eigenvector matrix.